

Outline

- Forecasting Revisited and the Solution of HW2
 1. Review of the HW2
- Optimize α and β .
 1. Numerical Simulation using Solver
 2. How to find the **best** tuning parameters practically?

[1] Forecasting Revisited and the Solution of HW2

Analytical Solution:

- Now, F is unknown.
- You need to use the estimated version of F , \hat{F} .

$$\underline{q^*} = \hat{F}^{-1}(TSL)$$

where

$$\underline{SL} = \frac{Cu}{Cu + Co} = \frac{r - c}{(r - c) + (c - s)} = \frac{r - c}{\underline{r - s}},$$

and \hat{F} is the CDF of demand **based on forecast**.

- Now, we are forecasting demand, so \hat{F} is not equal to F anymore.
- In the case of the additive demand, we can represent the solution as follows:

safety stock

$$\underline{q^*} = \underline{\text{Forecasted Demand}} + \underline{\text{Forecasting error}^{-1}(SL)}$$

Descriptive Analysis

- Prepare Rstudio and load general packages to use
- Load csv/txt file

```
df <- read.csv("../data/demand_data_session3.txt", sep = "\t")
```

```
# Add time trend
```

```
df <- df %>%  
  mutate(t = row_number()) %>%  
  as_tibble()
```

```
# Descriptives
```


```
summary(df)
```

##	t	demand
##	Min. : 1.0	Min. :1606
##	1st Qu.:125.8	1st Qu.:2477
##	Median :250.5	Median :3045
##	Mean :250.5	Mean :3012
##	3rd Qu.:375.2	3rd Qu.:3528
##	Max. :500.0	Max. :4410


Visualize

```
df %>%  
  ggplot(aes(x = demand)) +  
  geom_histogram()
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
df %>%  
  ggplot(aes(x = t, y = demand)) +  
  geom_line()
```



Setup

Set Exogenously Given Values as Constant

```
price <- 4
cost <- 0.8
salvage <- 0

# What we know
SL <- (price - cost) / (price - salvage) # 0.8
```

$$\frac{r-c}{r-s}$$

Set Blank Table

```
profit <- matrix(NA, nrow = nrow(df), ncol = 5)
colnames(profit) <- c("Simple Newsvendor", "Oracle", "Moving Average", "Simple ES", "Holt")
rownames(profit) <- paste0("day", 1:nrow(df))
head(profit)
```

##	Simple Newsvendor	Oracle	Moving Average	Simple ES	Holt
## day1	NA	NA	NA	NA	NA
## day2	NA	NA	NA	NA	NA
## day3	NA	NA	NA	NA	NA
## day4	NA	NA	NA	NA	NA
## day5	NA	NA	NA	NA	NA
## day6	NA	NA	NA	NA	NA

Simple Newsvendor

```
for (i in 301:nrow(df)) {  
  stocking_optimal_adaptive <- as.numeric(quantile(df$demand[1:(i-1)], SL))  
  profit[i, "Simple Newsvendor"] <-  
    price * min(stocking_optimal_adaptive, df$demand[i]) -  
    cost * stocking_optimal_adaptive  
}
```

Oracle

```
for (i in 301:nrow(df)) {  
  profit[i, "Oracle"] <- price * df$demand[i] - cost * df$demand[i]  
}
```



Moving average

```
df <- df %>%
  dplyr::mutate(forecast_MA4 =
    (dplyr::lag(demand, 4) +
     dplyr::lag(demand, 3) +
     dplyr::lag(demand, 2) +
     dplyr::lag(demand, 1)) / 4)

df <- df %>%
  dplyr::mutate(forecast_error_MA4 = demand - forecast_MA4)

for (i in 301:nrow(df)) {
  ✓ safety_stock <- as.numeric(quantile(df$forecast_error_MA4[1:(i-1)], SL, na.rm = TRUE))
  ✓ stocking_optimal_adaptive <- df$forecast_MA4[i] + safety_stock
  profit[i, "Moving Average"] <-
    {
      price * min(stocking_optimal_adaptive, df$demand[i]) -
      cost * stocking_optimal_adaptive
    }
}
```

[2] Optimizing tuning parameters

- Optimizing forecast based on the previous data (say, day 101-300)
- Use the optimized α to calculate the optimal stocking and average profit.
- Compare it with the naive one where we fixed α .

$$\begin{array}{r} 1 \quad -1 \quad / \\ 1 \quad +1 \quad / \\ \hline 2 \quad \quad 0 \end{array}$$

MSE

$$\frac{1}{n} \sqrt{\sum (x_i - \hat{x}_i)^2}$$

Demand forecasted

Simple Exponential Smoothing (untuned)

```
✓ L0 <- 2000
✓ alpha <- 0.2
Lt <- c(); Ft <- c()
for (i in 1:nrow(df)) {
  if (i == 1) {
    Lt[i] <- alpha * df$demand[i] + (1 - alpha) * L0
    Ft[i] <- L0
  } else {
    Lt[i] <- alpha * df$demand[i] + (1 - alpha) * Lt[i-1]
    Ft[i] <- Lt[i-1]
  }
}
df <- df %>%
  dplyr::mutate(forecast_Simple_ES = Ft,
               forecast_error_Simple_ES = demand - forecast_Simple_ES)

for (i in 301:nrow(df)) {
  ✓ safety_stock <- as.numeric(quantile(df$forecast_error_Simple_ES[1:(i-1)], SL, na.rm = TRUE))
  ✓ stocking_optimal_adaptive <- df$forecast_Simple_ES[i] + safety_stock
  ✓ profit[i, "Simple ES"] <-
    price * min(stocking_optimal_adaptive, df$demand[i]) -
    cost * stocking_optimal_adaptive
}
```

```
# Compare average profit by forecasting method  
apply(profit[301:500, ], 2, mean)
```

```
## Simple Newsvendor  
##      10583.16  
##           Holt  
##           NA
```

Oracle
11620.86

Moving Average
11272.74

Simple ES
11298.89

Construct MSE function for Simple Exponential Smoothing

```
mse_simple_ES <- function(alpha, predict_from, predict_to, df) {  
  
  df_predict <- df %>%  
    dplyr::slice(predict_from:predict_to)  
  df_true <- df %>%  
    dplyr::slice(predict_from:predict_to)  
  
  L0 <- 2000  
  Lt <- c(); Ft <- c()  
  for (i in 1:(predict_to - predict_from + 1)) {  
    if (i == 1) {  
      Lt[i] <- alpha * df_predict$demand[i] + (1 - alpha) * L0  
      Ft[i] <- L0  
    } else {  
      Lt[i] <- alpha * df_predict$demand[i] + (1 - alpha) * Lt[i-1]  
      Ft[i] <- Lt[i-1]  
    }  
  }  
  
  mse <- sum((df_true$demand - Ft)^2) / (predict_to - predict_from + 1)  
  return(mse)  
}
```

Optimize alpha

- There are many ways/solvers to do this.
- To impose the restriction $\alpha \in [0, 1]$, we used L-BFGS-B here.
- Other solver such as **BFGS**, **Nelder-Mead** should work well even without any range restriction in this simple data.

```
res <- optim(0.2, mse_simple_ES, method = "L-BFGS-B",  
            predict_from = 1,  
            predict_to = 300,  
            df = df,  
            lower = c(0),  
            upper = c(1))  
res
```

```
## $par  
## [1] 0.08900983  
##  
## $value  
## [1] 94325.41  
##  
## $counts  
## function gradient  
##      49      49  
##  
## $convergence  
## [1] 52  
##  
## $message  
## [1] "ERROR: ABNORMAL_TERMINATION_IN_LNSRCH"
```

Use the tuned values to calculate the optimal profit

```
profit <- cbind(profit, rep(NA, nrow(df)))  
colnames(profit)[6] <- "Simple ES (tuned)"
```

✓ `alpha <- res$par`

```
L0 <- 2000  
Lt <- c(); Ft <- c()  
for (i in 1:nrow(df)) {  
  if (i == 1) {  
    Lt[i] <- alpha * df$demand[i] + (1 - alpha) * L0  
    Ft[i] <- L0  
  } else {  
    Lt[i] <- alpha * df$demand[i] + (1 - alpha) * Lt[i-1]  
    Ft[i] <- Lt[i-1]  
  }  
}  
df <- df %>%  
  dplyr::mutate(forecast_Simple_ES = Ft,  
               forecast_error_Simple_ES = demand - forecast_Simple_ES)
```

```
for (i in 301:nrow(df)) {
```

✓ `safety_stock <- as.numeric(quantile(df$forecast_error_Simple_ES[1:(i-1)], SL, na.rm = TRUE))`

✓ `stocking_optimal_adaptive <- df$forecast_Simple_ES[i] + safety_stock`

✓ `profit[i, "Simple ES (tuned)"] <-
 price * min(stocking_optimal_adaptive, df$demand[i]) -
 cost * stocking_optimal_adaptive
}`

Comparison Table

```
# Compare average profit by forecating method  
apply(profit[301:500, ], 2, mean)
```

## Simple Newsvendor	Oracle	Moving Average	Simple ES
## 10583.16	11620.86	11272.74	<u>11298.89</u>
## Holt Simple ES (tuned)			
## NA	<u>11311.68</u>		

Solution: Holt's

```
# Load data
df <- read.csv("../data/demand_data_session2.txt")

# Add time trend
df <- df %>%
  mutate(t = row_number()) %>%
  as_tibble()

# Setting
price <- 4
cost <- 0.8
salvage <- 0

# What we know
optimal_order_quantity <- as.integer(quantile(df$demand, (price - cost) / (price - salvage)))

# Make table for summary
res_table <- data.frame(
  method = c("Oracle",
             "holt_wo_tuning",
             "holt_w_tuning"),
  AvgProf = rep(NA, 3)
)
```

SL

Oracle

```
# Oracle  
res_table$AvgProf[1] <- (price - cost) * mean(df$demand[301:500])
```

1. Estimate initial values

To ✓
✓
L,

```
res_regression <- lm(data = filter(df, t <= 300),  
                    formula = demand ~ t)  
b <- as.numeric(res_regression$coefficients[1])  
a <- as.numeric(res_regression$coefficients[2])
```

2 Forecast by Holt's method

```
✓ L_0 <- b
```

```
✓ T_0 <- a
```

```
{ alpha <- 0.2
```

```
{ beta <- 0.2
```

```
# use function make code clean, and easy to debug
```

```
forecast_holt_trend <- function(alpha, beta, starting_forecast, L_0, T_0, demand,  
                                forecast_start_day, forecast_end_day) {
```

```
  forecasted <- c()
```

```
  L_t <- c()
```

```
  T_t <- c()
```

```
  for (i in 1:(forecast_end_day-forecast_start_day)) {
```

```
    if (i == 1){
```

```
      # ex-ante
```

```
      forecasted[1] <- starting_forecast
```

```
      # ex-post (after observing the demand)
```

```
      L_t[1] <- alpha * demand[forecast_start_day+i] + (1 - alpha) * forecasted[1]
```

```
      T_t[1] <- beta * (L_t[1] - L_0) + (1 - beta) * T_0
```

```
    } else {
```

```
      # ex-ante: update the forecast with alpha and beta
```

```
      forecasted[i] <- L_t[i-1] + T_t[i-1]
```

```
      # ex-post: update the L_t and T_t
```

```
      L_t[i] <- alpha * demand[forecast_start_day+i] + (1 - alpha) * forecasted[i] #L_t[i-1]
```

```
      T_t[i] <- beta * (L_t[i] - L_t[i-1]) + (1 - beta) * T_t[i-1]
```

```
    }
```

```
  }
```

```
  output <- list(forecasted = forecasted,
```

```
                L_t = L_t,
```

```
                T_t = T_t)
```

```
  return(output)
```

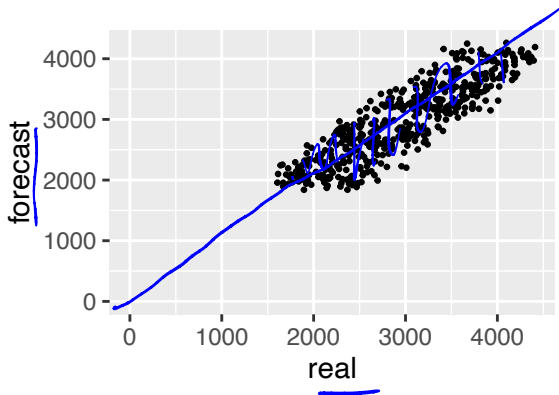
```
}
```

```

forecast_1_500 <- forecast_holt_trend(alpha = alpha,
                                     beta = beta,
                                     starting_forecast = b + 1 * a,
                                     L_0 = L_0,
                                     T_0 = T_0,
                                     demand = df$demand,
                                     forecast_start_day = 0,
                                     forecast_end_day = 500)

tibble(real = df$demand, forecast = forecast_1_500$forecasted) %>%
  ggplot(aes(x = real, y = forecast)) +
  geom_point(size = 0.5) +
  xlim(c(0, NA)) +
  ylim(c(0, NA))

```



Check: Example 7.3

- Observed demand (in thousands) has been
 $\left\{ \begin{array}{l} D_1 = 8,415, D_2 = 8,732, D_3 = 9,014, D_4 = \\ 9,808, D_5 = 10,413, \text{ and } D_6 = 11,961. \end{array} \right.$
- Forecast demand for Period 7 using trend-corrected exponential smoothing with $\alpha = 0.1, \beta = 0.2$.
- The first step is to obtain initial estimates of level and trend using linear regression.

F_7

- For the MP3 player data, we obtain

$$L_0 = 7,367 \text{ and } T_0 = 673$$

- The forecast for Period 1 is thus given by

$$\rightarrow F_1 = L_0 + T_0 = 7,367 + 673 = 8,040$$

- The observed demand for Period 1 is $D_1 = 8,415$.
- So, the error for Period 1 is thus given by

- With $\alpha = 0.1, \beta = 0.2$, the revised estimate of level and trend for Period 1:

$$\begin{aligned}\checkmark L_1 &= \alpha D_1 + (1 - \alpha)(L_0 + T_0) \\ &= (0.1 \times 8,415) + (0.9 \times 8,040) = 8,078\end{aligned}$$

$$\begin{aligned}\checkmark T_1 &= \beta(L_1 - L_0) + (1 - \beta)T_0 \\ T_1 &= [0.2 \times (8,078 - 7,367)] + (0.8 \times 673) = 681\end{aligned}$$

- We thus obtain the following forecast for Period 2:

$$\underline{F_2} = L_1 + T_1 = 8,078 + 681 = 8,759$$

- Continuing in this manner, we obtain

$$\begin{aligned}&\bullet L_2 = 8,755, T_2 = 680, L_3 = 9,393, T_3 = 672, \\ &L_4 = 10,039, T_4 = 666, L_5 = 10,676, T_5 = 661, L_6 = \\ &11,399, T_6 = 673.\end{aligned}$$

- This gives us a forecast for Period 7 of

$$\underline{F_7} = L_6 + T_6 = 11,399 + 673 = \underline{12,072} \quad \checkmark$$

Debug if any: Check the Match!


```
demand_ex <- c(8415, 8732, 9014, 9808, 10413, 11961)
alpha_ex <- 0.1
beta_ex <- 0.2

L_0_ex <- 7367; T_0_ex <- 673
F1 <- L_0_ex + T_0_ex

forecast_ex <- forecast_holt_trend(alpha = alpha_ex,
                                   beta = beta_ex,
                                   # corresponds to F1 above
                                   starting_forecast = L_0_ex + 1 * T_0_ex,
                                   L_0 = L_0_ex,
                                   T_0 = T_0_ex,
                                   demand = demand_ex,
                                   forecast_start_day = 0,
                                   forecast_end_day = 7)

tibble(t = 1:7, forecast = forecast_ex$forecasted)
```

```
## # A tibble: 7 x 2
##       t forecast
##   <int>   <dbl>
## 1     1    8040
## 2     2    8758
## 3     3    9435.
## 4     4   10065.
## 5     5   10706.
## 6     6   11337.
## 7     7   12072.
```



No bugs, you are good to go.

3. Optimal Stocking at week 301

```
✓ forecast_error_holt <- df$demand - forecast_1_500$forecasted  
✓ safety_stock <- quantile(forecast_error_holt[1:300], 0.80)  
✓ stocking <- forecast_1_500$forecasted[300] + safety_stock
```

SL

$$\frac{r-c}{r-s}$$

4. Ave.Profit in Holt's method?

```
✓ forecast_error_holt <- df$demand - forecast_1_500$forecasted
✓ safety_stock <- quantile(forecast_error_holt[1:300], 0.80)
✓ stocking <- forecast_1_500$forecasted[301:500] + safety_stock
```

```
# Ave. profit
res_table[2, "AvgProf"] <-
  mean(price * pmin(stocking, df$demand[301:500]) +
        salvage * pmax(stocking - df$demand[301:500], 0) -
        cost * stocking)
res_table
```

General Formula

```
##          method AvgProf
## 1      Oracle 11620.86
## 2 holt_wo_tuning 11283.66
## 3 holt_w_tuning      NA
```

5. optimization

```
# Let's find optimal alpha and beta
mse_holt_2 <- function(par, starting_forecast, L_0, T_0, demand,
  forecast_start_day, forecast_end_day) {
  alpha <- par[1]
  beta <- par[2]

  forecasted <- forecast_holt_trend(alpha = alpha,
    beta = beta,
    starting_forecast = starting_forecast, #b + 100 * a,
    L_0 = L_0, #b + 100*a,
    T_0 = T_0, #a,
    demand = demand,
    forecast_start_day = forecast_start_day,
    forecast_end_day = forecast_end_day)[["forecasted"]]

  ✓ mse <- sqrt(sum((demand[(forecast_start_day+1):(forecast_end_day)] - forecasted)^2) / length(forecasted))
  return(mse)
}
```

```
res <- optim(par = c(0.2, 0.2), mse_holt_2, method = "L-BFGS-B",  
  lower = c(0, 0),  
  upper = c(1, 1),  
  starting_forecast = b + 1 * a,  
  L_0 = b + 1*a,  
  T_0 = a,  
  demand = df$demand,  
  forecast_start_day = 0,  
  forecast_end_day = 300)  
res
```

```
## $par  
## [1] 0 0 ✓  
##  
## $value  
## [1] 294.4498  
##  
## $counts  
## function gradient  
##      2      2  
##  
## $convergence  
## [1] 0  
##  
## $message  
## [1] "CONVERGENCE: NORM OF PROJECTED GRADIENT <= PGTOL"  
alpha <- res$par[1]  
beta <- res$par[2]
```

Alternative: Grid-Search

```
# You can also do the grid-search
tb_mse <- expand_grid(alpha = seq(from = 0, to = 1, length.out = 100),
                      beta = seq(from = 0, to = 1, length.out = 100)) %>%
  rowwise() %>%
  mutate(mse = mse_holt_2(par = c(alpha, beta),
                           starting_forecast = b + 1 * a,
                           L_0 = b + 1*a,
                           T_0 = a,
                           demand = df$demand,
                           forecast_start_day = 0,
                           forecast_end_day = 300))

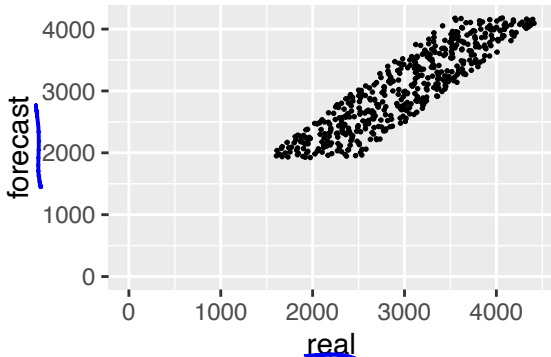
grid_1 <- unique(tb_mse$alpha)
grid_2 <- unique(tb_mse$beta)
obj_spread <- matrix(0, nrow = length(grid_1), ncol = length(grid_2))
for (i in 1:length(grid_1)) {
  for (j in 1:length(grid_2)) {
    temp <- tb_mse %>%
      filter(alpha == grid_1[i] & beta == grid_2[j])
    obj_spread[i,j] <- temp[["mse"]]
  }
}
saveRDS(obj_spread, file = "../code/grid_search.rds")

obj_spread <- readRDS("../code/grid_search.rds")
library(plotly)
library(htmlwidgets)
p <- plot_ly(z = obj_spread, type = "surface")
htmlwidgets::saveWidget(p, "../code/obj_spread.html", selfcontained = F, libdir = "lib")
```

6. Use the tuned parameters to calculate Ave Profit

```
forecast_1_500 <- forecast_holt_trend(alpha = alpha,
                                     beta = beta,
                                     starting_forecast = b + 1 * a,
                                     L_0 = L_0,
                                     T_0 = T_0,
                                     demand = df$demand,
                                     forecast_start_day = 0,
                                     forecast_end_day = 500)

tibble(real = df$demand, forecast = forecast_1_500$forecasted) %>%
  ggplot(aes(x = real, y = forecast)) +
  geom_point(size = 0.3) +
  xlim(c(0, NA)) +
  ylim(c(0, NA))
```



```
forecast_error_holt <- df$demand - forecast_1_500$forecasted
safety_stock <- quantile(forecast_error_holt[1:300], 0.80)
stocking <- forecast_1_500$forecasted[301:500] + safety_stock
```

```
# Ave. profit
res_table[3, "AvgProf"] <-
  mean(price * pmin(stocking, df$demand[301:500]) +
        salvage * pmax(stocking - df$demand[301:500], 0) -
        cost * stocking)
res_table
```

```
##           method AvgProf
## 1         Oracle 11620.86
## 2 holt_wo_tuning 11283.66
## 3 holt_w_tuning 11288.88
```

Some common Mistakes

- Mistake in forecasting
 - Inconsistent initial values, setup, parameters.
- Mistake in MSE
 - Using ES's MSE, which is wrong.
 - Some bugs inside MSE function.
- Mistake in optimal stocking
 - You cannot use the future value to decide optimal stocking.
- Mistake in optimization
 - Not optimizing two parameters.
 - Not feeding appropriate local objects.

I highly recommend that you review HWs as this might share the large portion in final exam too.

Comment

- In grading, we allow you to have different values if that is not tremendously different.
- Those “small” differences might come from the following difference.
 - Adaptive or Fixing error distribution?
 - Slight difference in tuned parameters.
 - The way to initialize the Holt’s method.
- Debug!
 - Check if you get the “exact” same result when you apply your function to some small scale setup.